

Spam Classification & Spam Filtering

Girish Sastry

May 7, 2011

Abstract

Spam is a concept that has been around a long time. These unsolicited messages dilute the message pool, causing frustration and leading to unhappy consumers. With the advent of the internet and electronic mail, the barrier to entry for spamming has become very low. Recipients of spam e-mail messages must waste time and effort to delete irrelevant spam e-mails. This paper focuses on the analysis of different methods for spam classification. We compare the learning performance of different types of classifiers on the SpamBase e-mail data set. We also look at different tools and programming languages to express these machine learning computations. We conclude by discussing the tradeoffs of our different methods and our tools, and point to the results we learned by applying these machine learning methods to the SpamBase data set.

1 Introduction

1.1 Spam

In today's digital age, concern about the reception of unsolicited junk e-mail, ("Spam"), has been increasing. When spam e-mail is received in small quantities, it may not be a hassle for users to manually delete unwanted e-mails. However, because the barrier of entry for spammers is so low for e-mail, users would have to spend a lot of time and resources to keep their inboxes clean and void of spam. For many internet companies whose business models are based on advertising, spam is the invisible nemesis that sucks away advertising money. For example, companies such as Facebook, Google, and Twitter, all make more money based on advertisement impressions. The more time a user spends on these sites, the more likely they will be to click an ad. The presence of spam reduces this and presents many data curation issues. Thus, spam filtering is an important problem in today's internet age.

How do we recognize a spam e-mail? There are many features present in the body of the e-mail message that may indicate spam. In a 1998 study, authors reported the different purposes of reported spam e-mails[1]:

Similar to the telemarketers of old, spam e-mails want to sell stuff to users. Recognizing this from the features present in the body of an e-mail is hard, as these features are probably not globally constant – they would vary not only from spammer to spammer, but from user to user. One man's gold may be another man's trash. Thus, in practice, spam filtering must be robust globally, but also personalized to a user's specific needs.

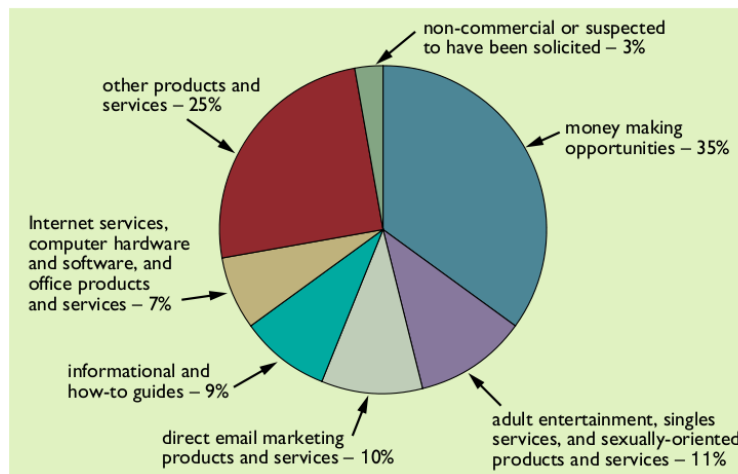


Figure 1: Type of products advertised by spam e-mails

1.2 SpamBase Data Set

The SpamBase data set, available for free from the UCI Machine Learning repository, was used for spam classification for this project. The SpamBase data set was collected internally by Hewlett Packard Labs for research purposes and then released to the public. The spam emails were self-reported as being spam. While spam for one e-mail user may not be spam for another, the data set attempts to extract some relevant features in order to make this a machine learning problem. There are fifty-five features total. Forty-eight of these features are continuous, real-valued attributes for word frequencies. Six are continuous, real-valued attributes for character frequencies. The rest of the features take into account other commonly seen features in spam: excessively long capitalized strings. Figure 2 shows a scatterplot matrix of the first five features in the data set. It is important to note that due to the nature of the data set, we will not see strong correlations immediately present. There is a wide variance in the features for spam emails – once again, one person’s spam is another person’s coupon.

1.3 Goals

For this project, we have two overall goals.

1.3.1 Classification Performance

First, we would like to measure and compare the classification performance of a number of methods. These include methods discussed in class and out of class. For this project, we consider:

- Decision Trees
- k-Nearest Neighbor

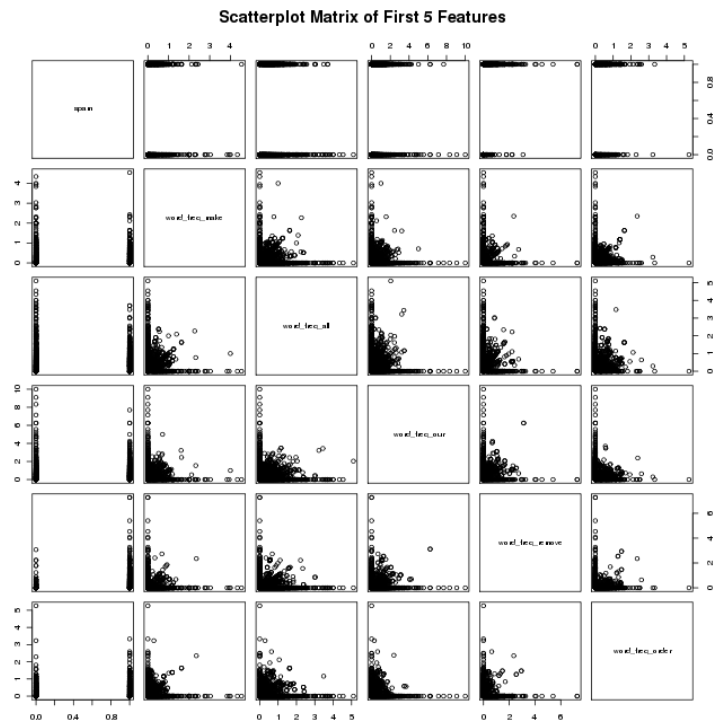


Figure 2: Scatter plot matrix of first 5 features from SpamBase

We would like to minimize both misclassification error and false positives. In the case of spam, false positives very bad.

1.3.2 Tools to Express Computations

Next, we would like to explore the different tools available to express these computations. For most of the course we have been using R, the “statistical programming language.” However, there are some limitations of R that may not make it the best tool for the task. For example, R’s data frame’s are stored in memory. This means that R is not suited for large data sets that do not fit in memory. With today’s petabytes and petabytes of data being produced on the web, it is time to consider other tools. For this project, we use the following programming languages and toolkits:

- WEKA Toolkit (Java)
- Clojure
- R

Clojure is a dialect of Lisp, and has recently exploded in popularity. There have been several machine learning packages written in Clojure. The appeal of Clojure as a general purpose programming language is quite clear. It is a Lisp, so we obtain all the benefits of a dynamically typed, functional programming

language with minimal syntax. This means that computations can be expressed more succinctly than canonical imperative programming languages like C or Java. Functions are passed around as values, so, just as in mathematics, functions can be composed and passed as arguments to functions.

For statistical computing, Lisp has recently drawn speculation as the next big thing, and as R's successor[4]. R's limitations include:

- Pass-by-value computations: only local values of parameters are modified
- Scalar operations: Vectorized operations execute much faster than scalar operations, but are harder to implement
- Native code compilation: R is slow. Thus, in many intensive use cases, the performance-critical code is moved to C.
- Rudimentary type system: Type checking, inference, and analysis is a real workhorse of the compiler, and can aid the programmer in debugging. R lacks this. [4]

Clojure is a Lisp that operates on the Java Virtual Machine, which means that it has access to the huge assortment of Java libraries. This makes it an ideal candidate for a new “tool” to experiment with in this project.

2 Results & Discussion

2.1 k-Nearest Neighbors

The k-Nearest Neighbors classification algorithm is one of the simplest classification algorithms. Indeed, it is the first algorithm, other than univariate regression, that we covered in this course. k-Nearest Neighbors attempts to classify an object by its nearest neighbors in the feature space. An object is classified by a majority vote of its neighbors. If k is small, say 1, then an object “takes on” the properties of its nearest neighbor.

The distance metric used to find the “nearest” neighbor is usually the Euclidean distance metric, and this is what is used in our experiments. For some classification problems in information retrieval and text classification, other distance metrics, such as Hamming Distance, may be used[3].

It is up to the experimenter to set the value of k in the k-Nearest Neighbors algorithm. There are a couple things to consider when setting this value. An extremely low value of k is extremely susceptible to noise or irrelevant, less important features in the data set. Intuitively, for the SpamBase data set, this could be something like `char_freq_()`, where it is not immediately clear what sort of correlation this could have with spam or non-spam emails. We go through 10-fold cross validation to choose the optimal value of k for SpamBase classification – in our case, $k = 5$.

One drawback to the majority voting that is present in the k-Nearest Neighbors algorithm is that classes that tend to appear more frequently in the data set (i.e. more instances) also tend to dominate the classification results. In our case, more than 50% of the SpamBase data set is “non-spam” e-mails, so we may see this domination.

2.2 Naive Bayes

Next, we run a Naive Bayesian classifier on the SpamBase data set. Naive bayes classifiers are computationally very fast, relatively easy to implement, and relatively easy to understand. In addition, and perhaps in spite of their simplicity, Naive Bayes classifiers have worked pretty well in practice[6]. These attributes make Naive Bayes classifiers a prime classifier choice for this set of experiments.

Naive bayes classifiers are simple probabilistic based on Bayes' Theorem. The "naive" in Naive Bayes refers to the strong independence assumption inherent in these classifiers – each feature is assumed to be independent of other features. Assuming n features ranging from F_1, F_2, \dots, F_n , and a class variable called C , there are a few inputs to the probability model:

- Prior probabilities: $P(C)$
- Likelihood: $P(F_1, \dots, F_n|C)$
- Evidence: $P(F_1, \dots, F_n)$

And we can write, using Bayes Theorem, the formula to calculate the posterior probability:

$$P(C|F_1, \dots, F_n) = \frac{P(C)P(F_1, \dots, F_n|C)}{P(F_1 \dots F_n)}$$

In practice, the "evidence" probability is usually a constant, since it doesn't depend on the class variable, C .

There are a couple reasons why the assumption of independence is an ok assumption. First, we just need the relative probabilities in order to make an accurate prediction. These relative probabilities just need to be ranked correctly. We do not need the extra information regarding class probabilities. This implies that Naive Bayes is a pretty robust classification method. Second, variable decoupling has some positive side-effects. Variable decoupling reduces higher order interaction terms that take away from the predictive model and thus helps reduce the curse of dimensionality. Thus, the "naive" assumption of strong independence that is present in the model is not as naive as it initially seems.

The overall classification model is as follows[7]:

$$\text{classify}(f_1, \dots, f_n) = \underset{c}{\operatorname{argmax}} p(C = c) \prod_{i=1}^n p(F_i = f_i|C = c).$$

Figure 3: Classification model for Naive Bayes

We obtain a training error of 10.60% and a test error of 10.43% using the Naive Bayes classifier. This misclassification error is not that good. It appears that the assumptions that Naive Bayes makes, the very ones that we argued were ok assumptions to make, skew the results. We also tried two other Naive Bayes algorithms from the WEKA machine learning toolkit and obtained similar results. `ComplementNaiveBayes` provided a slightly better result of 9.14% training error and 9.58% test error. While better, this is still not the best error we could get.

2.3 Support Vector Machines

Support vector machine is a an interesting learning method that is a non-probabilistic binary classifier. They work by predicting for each instance of data the class that the data is a part of.

The model that a support vector machine builds works as follows: the support vector machine constructs a hyperplane (or a set of hyperplanes for a higher dimensional space). This hyperplane that is drawn can be described by the characteristic equation:

$$\mathbf{w}\mathbf{x} + b = 0$$

where

- \mathbf{w} is normal to the hyperplane,
- $\frac{b}{\|\mathbf{w}\|}$ is the distance from the origin to the hyperplane

The term “support vector” refers to the vectors, or instances of data, that are closest to the hyperplane. The goal of support vector machines is to orient the hyperplane so that the support vectors are as far as possible from the hyperplane.

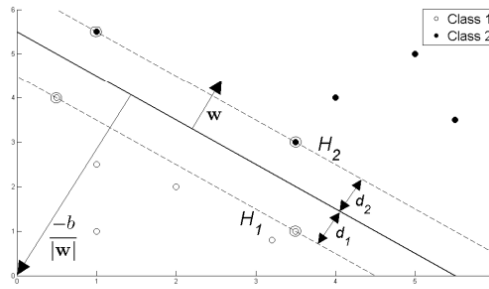


Figure 1: *Hyperplane through two linearly separable classes*

Figure 4: Support vectors & hyperplane for two classes, example

Thus, support vector machines attempt to assign values to \mathbf{w} and b for the data set.

It should be noted that this theoretical model of support vector machines assumes that the data is linearly separable into two classes. To deal with data that is not fully linearly separable, the model relaxes some of the optimization constraints when choosing values for \mathbf{w} and b [2].

On the spam data set, we obtain a training error of 6.50% and a test error of 6.38%. This is a decent result, and beats the 7% error that is considered “good.” This better performance can be explained by some of the advantages of support vector machines. Support vector machines often perform better when the data set is small and noisy. This is the case with the SpamBase data set. We have a relatively small number of instances. The data set is noisy because of the nature of spam. One person may report something with `word_freq_hp` as spam, while it is not to another person.

2.4 Decision Trees

First, we consider decision trees. Decision trees are a simple procedure that classifies instances of data. They work by constructing a tree of decisions, where the leaf nodes represent classifications, and each intermediate node represents a decision based on a conjunction of features that leads to a classification.

Using unpruned decision trees, we obtain a training error of 10.53% and a test error of 9.98%.

Next, we turn an ensemble method of classification using decision trees. We use boosting on the SpamBase data set. Boosting is based on the simple question: “Can a set of weak learners create a single strong learner?” [5] Boosting works, in the general case, by iteratively learning weak classifiers and then weighting and adding them to a final strong classifier. While there are many algorithms for boosting, we use Adaboost here. Boosting seeks to minimize the convex loss function

$$\sum_i e^{-y_i f(x_i)}$$

while seeking a function

$$f = \sum_t \alpha_t h_t$$

On the SpamBase data set, boosting gives a training error of 5.00% and a test error of 5.15%. This is much better than any of the other methods so far. This is probably due to the ensemble method of classification. At a high level, AdaBoost minimizes the exponential loss function. Thus, it basically chooses a hyperplane and orients it away from the two classes, in a similar way to support vector machines. It seems that based on these results that ensemble methods are the way to go. To further test this, we look at support vector machines next.

Next, we turn to another ensemble method of classification using decision trees. Random Forests are a way of using a slightly perturbed version of the data set in simulation along with multiple decision trees to classify data. Random forests combine the idea of “bagging” and random feature selection to produce multiple decision trees under controlled variation. The learning algorithm is as follows[3]:

1. Init: Let N be the number of instances in training set, and the number of features be M
2. We know m input variables which are used to determine the decision at each node of the tree. $m \ll M$.
3. Choose a training set: Choose N times with replacement from all N training instances
4. For each node in the tree, randomly choose m variables used to make the decision at that point.
5. Let each tree grow and do not prune

We obtain a training set error of 4.79% and a test set error of 4.69%. The confusion matrices for both the training and test sets are reported as follows:

Our training and test error rates are very good. They beat the “good” error rates by almost half. The number of false positives did not vary that much from our other tree based results, but there are still false positives. It seems that random forests, and perhaps ensemble methods in general, do much better at predicting spam. Because of possible high inter-variable correlation that many of other methods do not account, random forests perform better.

3 Discussion

It is interesting to note that Naive Bayes did not provide us with good results. In the Computer Science field, most companies that use spam filters in production use Naive Bayes as the workhorse of their spam filters. It is evident, based on our results here, and intuition, that these same institutions also use ensemble methods for classification. An untuned Naive Bayes algorithm does not provide good results, but the ensemble methods of random forests and AdaBoost do. Thus, it makes sense that spam filters in production use some combination of learners on top Naive Bayes to implement their spam filters.

We experimented with different tools of Clojure, WEKA, and R in this project. We find that Clojure has probably the most potential to write real machine learning systems. It is a very well designed language with many features and operates on the JVM. However, the current landscape of machine learning libraries is very sparse. R just completely wins this category with easily obtainable packages from CRAN and other repositories. The syntax and semantic constructs of Clojure are very nice. It is very easy to compose your own functions and build up in a modular style.

For all methods used, forcing a value of zero false-positives raises the training and test errors to 22%-25%. It seems hard to balance the two factors of minimizing false-positives and misclassification error.

4 Conclusions

We conclude that ensemble methods for classification provide better results than singular methods for classification. We also conclude that real production spam filters are produced by some variation of ensemble methods. It seems that Naive Bayes alone is very bad, but perhaps combined with other methods it improves drastically.

References

- [1] Lorrie Faith Cranor and Brian A. LaMacchia. Spam. *Communications of The ACM*, 41:74–83, 1998.
- [2] Tristan Fletcher. Support vector machines explained.
- [3] Trevor Hastie, Robert Tibshirani, and Jerome Friedman. *The Elements of Statistical Learning*. Springer Series in Statistics. Springer New York Inc., New York, NY, USA, 2001.

- [4] Ross Ihaka and Duncan Temple Lang. Back to the future: Lisp as a base for a statistical computing system.
- [5] Michael Kearns. Thoughts on hypothesis boosting, 1998.
- [6] F Lotte, M Congedo, A Lécuyer, and F Lamarche. Topical review a review of classification algorithms for eeg-based brain-computer interfaces.
- [7] Wikipedia. Wikipedia article on naive bayes.

5 Appendix